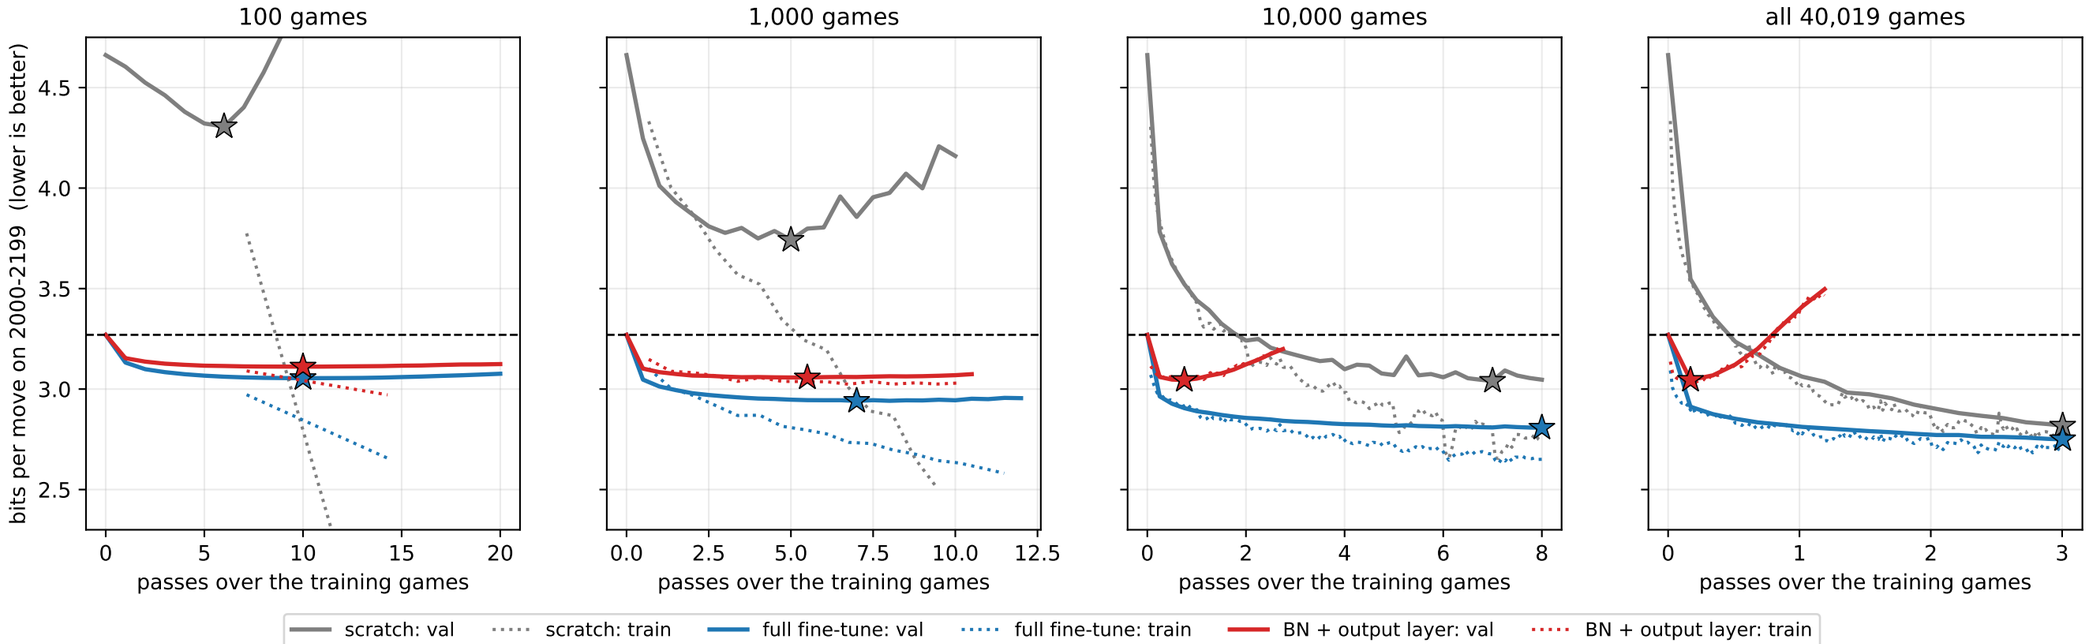

Seed 0. Solid = val bits on 42,433 held-out 2000-2199 positions; dotted = train loss (average of 50 steps); ★ = best checkpoint (the one tested). Black dashed = base model with 0 strong games (val 3.270). At 100 and 1k games the train loss keeps falling after the ★ while val rises (memorizing). BN + output layer: flat train loss at 100 and 1k games; at 10k and all games its train loss rises again after a few hundred steps, so the optimisation diverges (lr 1e-3), which says nothing about its capacity. With 10k and all games full fine-tuning and scratch were still improving when they hit the pass limit.