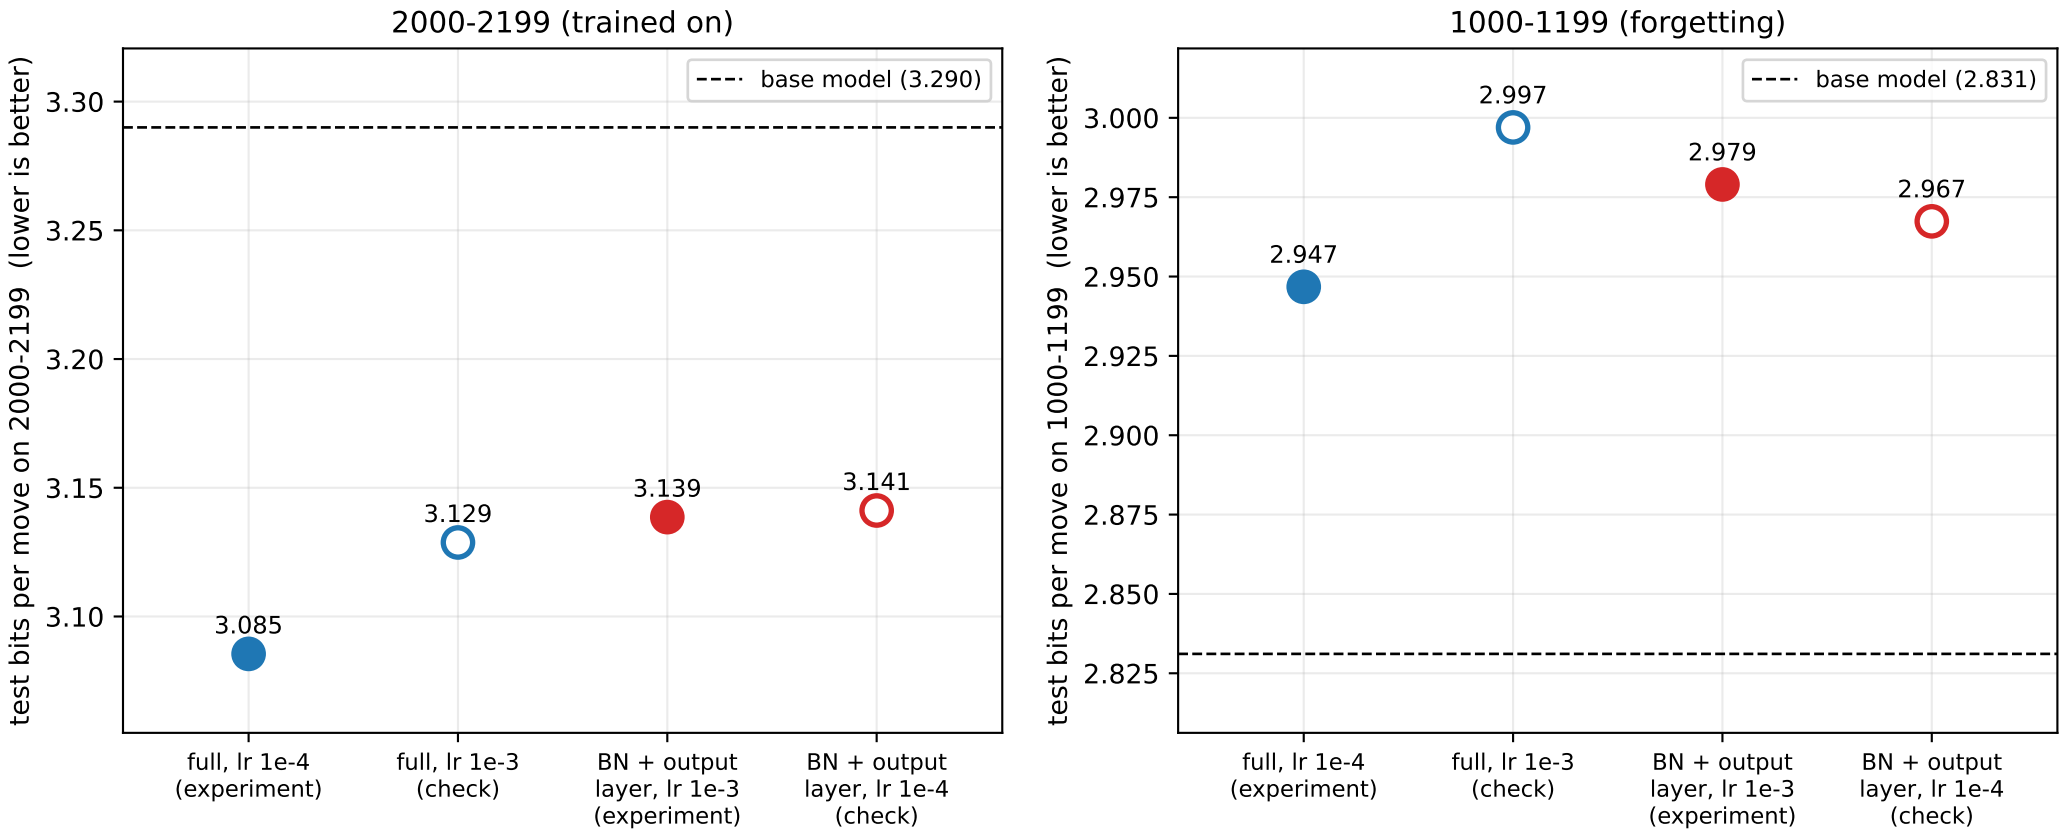

N = 100 games, seed 0 only (not checked at larger N). Filled = the learning rate used in the experiment; open = the other option's learning rate. BN + output layer at 1e-4 reaches 3.141 against 3.139 at 1e-3, after 6× more steps (and still improving at the step limit); full fine-tuning at 1e-3 is best after 2 passes and then overfits. So at N = 100 the full-vs-BN gap is not caused by the learning rate.