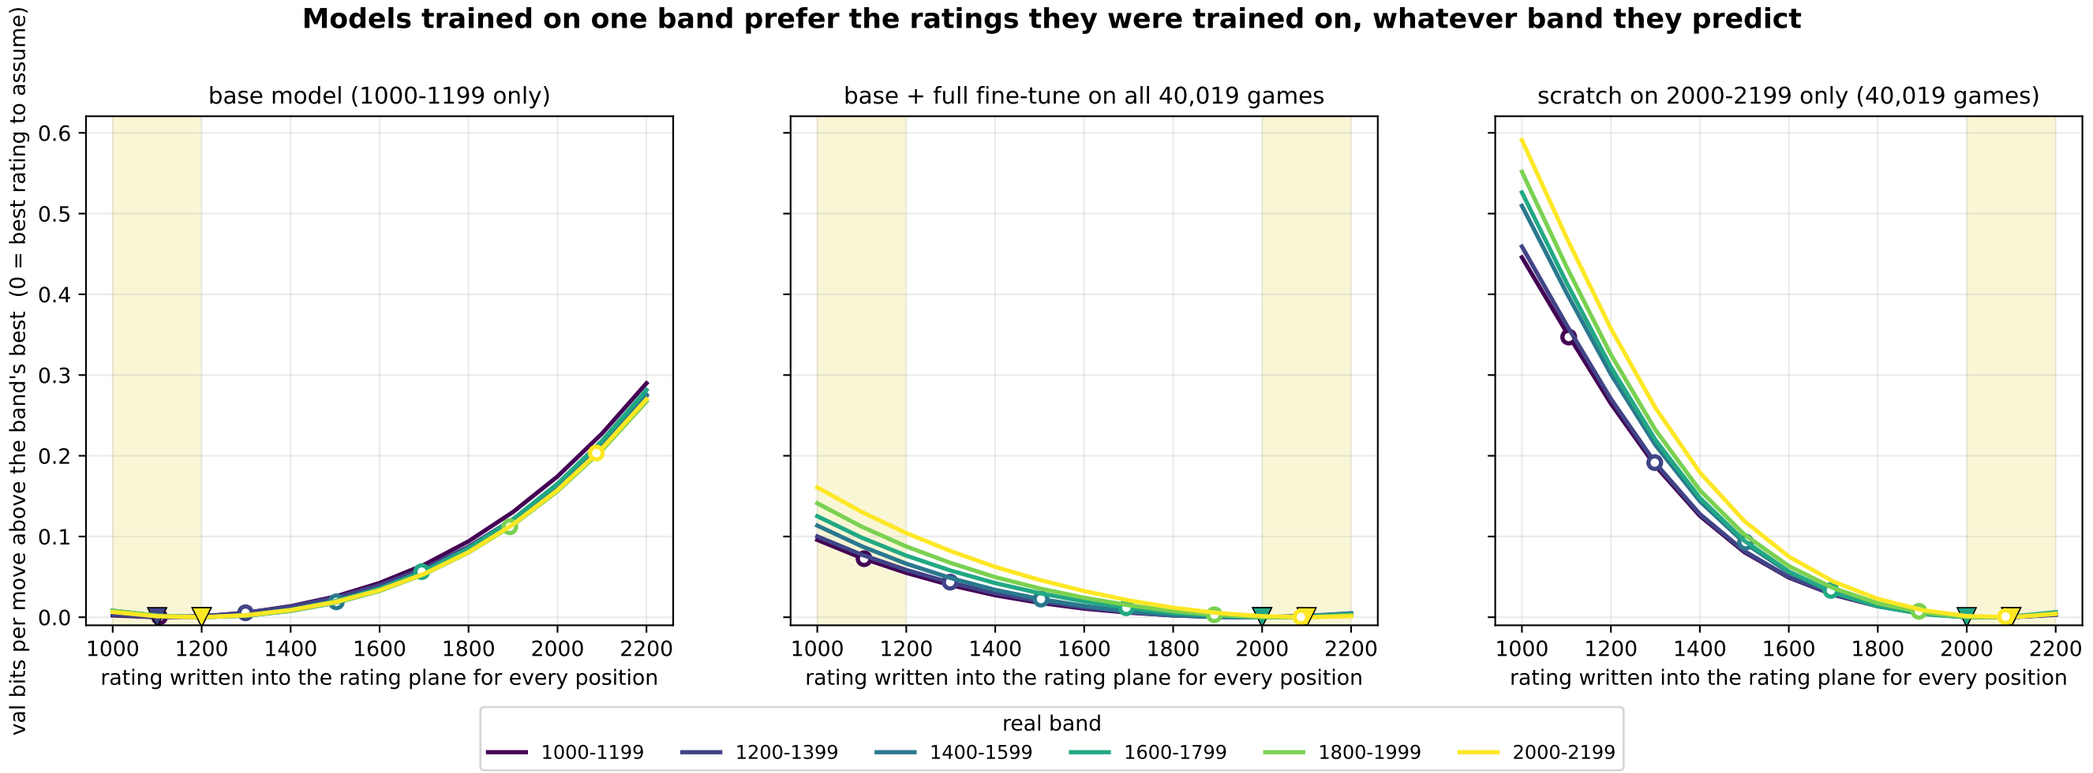


Each line is one band's val positions (shard 000), measured with the rating plane set to the same value for every position. ▼ = the rating that predicts that band best; ○ = the band's real ratings (measured normally). Shaded = ratings the model saw in training. A model that had learned what a rating means would predict each band best near its real ratings (▼ on ○). Instead, for every band, the base model prefers to be told ≈1100-1200 and the other two ≈2000-2100. The bottom of each curve is flat (best and second-best rating within 0.0020 bits), so only the location of the curve is meaningful, not the exact best rating. Telling the base model the real ratings of 2000-2199 players costs 0.20 bits compared with telling it 1200.